

000
001
002
003
004
005
006
007
008
009
010
011
012
013
014
015
016
017
018
019
020
021
022
023
024
025
026
027
028
029
030
031
032
033
034
035
036
037
038
039
040
041
042
043
044
045
046
047
048
049
050
051
052
053

Qualitative and Quantitative Evaluation of GANs on Pizza Dataset

Adithya Shrivastava
as3652
Rutgers University
Department of Computer Science
as3652@scarletmail.rutgers.edu

Janish Parikh
jrp328
Rutgers University
Department of Computer Science
janish.parikh@rutgers.edu

Nishant Dhargalkar
nsd78
Rutgers University
Department of Computer Science
nsd78@scarletmail.rutgers.edu

Jiawei Tang
jt1039
Rutgers University
Department of Computer Science
jiawei.tang@rutgers.edu

Abstract

Generative Adversarial Networks (GANs) are deep-learning based generative models. They can be used for a wide variety of purposes like style transfer, photo blending, image-to-image translation, etc. In this task, we aim to use GAN training real pizza and synthetic pizza datasets so that we can generate our own synthetic pizza.

1. Introduction

Generative adversarial networks (GANs) are a generative model with implicit density estimation, part of unsupervised learning and are using two neural networks. Thus, we understand the terms “generative” and “networks” in “generative adversarial networks”. This report displays our work for step 1.

2. Pre-processing

2.1. Datasets

The datasets included in this task are downloaded from <https://drive.google.com/drive/folders/1jaag8cyb72bFHSA-V8ckieTL9uAOsK7r?usp=sharing> [4]. We train models on two datasets: SyntheticPizza 2 and RealPizza datasets 2.1. The SyntheticPizza images and RealPizza images are resized to 256×256 pixels and normalized to $[-1, 1]$.



Figure 1. Real Pizza Images



Figure 2. Synthetic Pizza Images

054
055
056
057
058
059
060
061
062
063
064
065
066
067
068
069
070
071
072
073
074
075
076
077
078
079
080
081
082
083
084
085
086
087
088
089
090
091
092
093
094
095
096
097
098
099
100
101
102
103
104
105
106
107

108
109
110
111
112
113
114
115
116
117
118
119
120
121
122
123
124

3. Implementations

3.1. Generator architectures

We adopt the architecture from [3]. Use 9 residual blocks for 256×256 training images. Let $c7s1 - k$ denote a 7×7 Convolution-InstanceNorm-ReLU layer with k filters and stride 1. dk denotes a 3×3 Convolution-InstanceNorm-ReLU layer with k filters and stride 2. Reflection padding was used to reduce artifacts. Rk denotes a residual block that contains two 3×3 convolutional layers with the same number of filters on both layer. uk denotes a 3×3 fractional-strided-ConvolutionInstanceNorm-ReLU layer with k filters and stride $\frac{1}{2}$. The network with 9 residual blocks consists of: $c7s1 - 64, d128, d256, R256, R256, R256, R256, R256, R256, R256, u128, u64, c7s1 - 3$.

3.2. Discriminator architectures

We use 70×70 PatchGAN[2]. Let Ck denote a 4×4 Convolution-InstanceNorm-LeakyReLU layer with k filters and stride 2. After the last layer, we apply a convolution to produce a 1-dimensional output. We do not use InstanceNorm for the first $C64$ layer. We use leaky ReLUs with a slope of 0.2. The discriminator architecture is:

$C64 - C128 - C256 - C512$

3.3. Training

Repeat step 3.3.1 and 3.3.2. The training of a GAN can be understood better from the below image.

Here we are generating a random noise z .

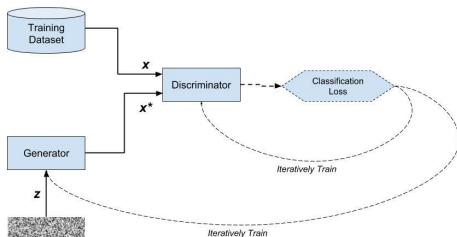


Figure 3. Generic GAN Training

It acts as the inout to our Generator network. This noise gets transformed into fake images with the help of the Generator.

Our dataset is either the real pizza images or the synthetic pizza images.

Discriminator and Generator are both Neural Networks that are trained together.

In GAN, a Discriminator competes with a Generator same as two players of zero-sum games.

3.3.1 Training Discriminator

Get real images from datasets and input them to the discriminator. Label the output as 1. Generate random noise z and input it to the generator to get $G(z)$, and then input $G(z)$ to the discriminator. Label the output as 0.

3.3.2 Training Generator

Generate random noise z and input it to the generator to get $G(z)$, and then input $G(z)$ to the discriminator. Label the output as 1.

3.4. Methodology

We used script files to train the GAN and persist the weights. Due to the rigorous disconnects and failures faced in using Jupyter Hub, we decide to move to ilabs. We executed python scripts in the ssh environment of ilabs and persisted the weights of generators and discriminators. However, we later realized that we have to submit completely executed colab notebooks. Hence, due to the shortage of time and carelessness from our end we could run lesser number of epochs to train our GAN in the Colab Notebook. We have shared both the colab notebooks and the python scripts used by us for the training of the GAN. I hope you understand.

4. Qualitatively Evaluation

4.1. Real Pizza Dataset

We ran 100 epochs in the initial setup and here are the images generated by the GAN at different epochs:



Figure 4. Output of GAN trained on Real Pizza dataset - Epoch 67



Figure 5. Output of GAN trained on Real Pizza dataset - Epoch 74

In the initial epochs the Generator learns the borders and structure of the pizza. In further epochs we can observe that the generator tries to learn the color of the pizza top (the sauces) and then toppings on the pizza. We can see that the generated pizza images are blurry in the toppings region.



Figure 6. Output of GAN trained on Real Pizza dataset - Epoch 87



Figure 7. Output of GAN trained on Real Pizza dataset - Epoch 89

4.2. Synthetic Pizza Dataset

We ran 50 epochs in the initial setup and here are the images generated by the GAN at different epochs:

In the initial epochs the Generator trained on the Syntehtic



Figure 8. Output of GAN trained on Synthetic Pizza dataset - Epoch 30

Pizza Dataset learns the borders and structure of the pizza. In further epochs we can observe that the generator tries to learn the sauce and cheese on the pizza. We can see that the generated synthetic pizza images are blurry as the number of epochs it is trained for is less.

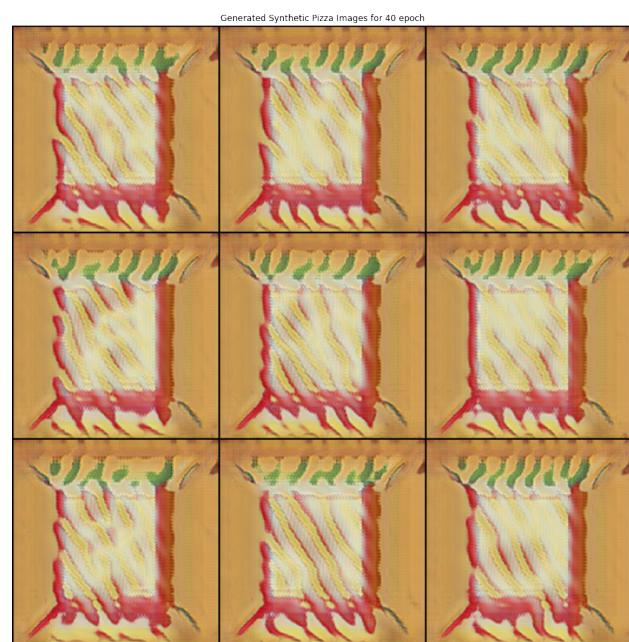


Figure 9. Output of GAN trained on Synthetic Pizza dataset - Epoch 40

4.3. Metrics

Real Pizza Dataset

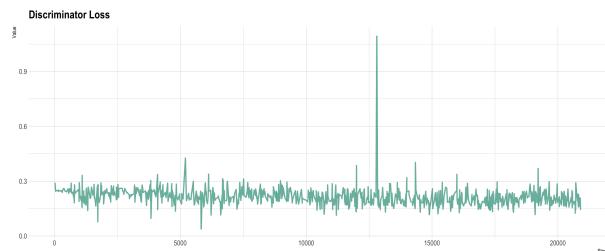


Figure 10. Discriminator Loss

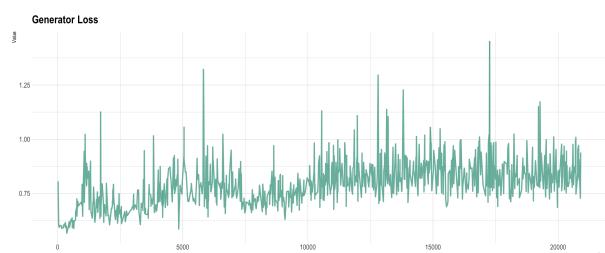
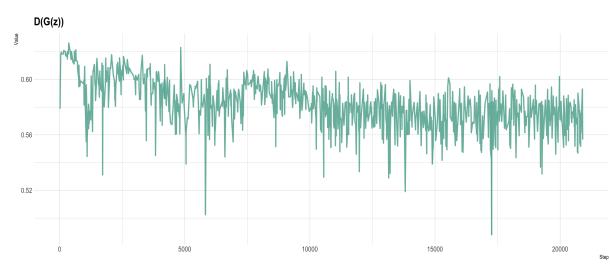
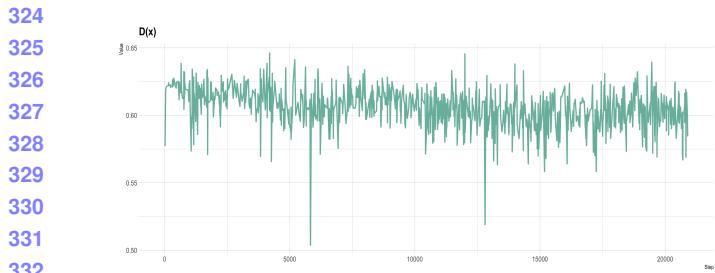
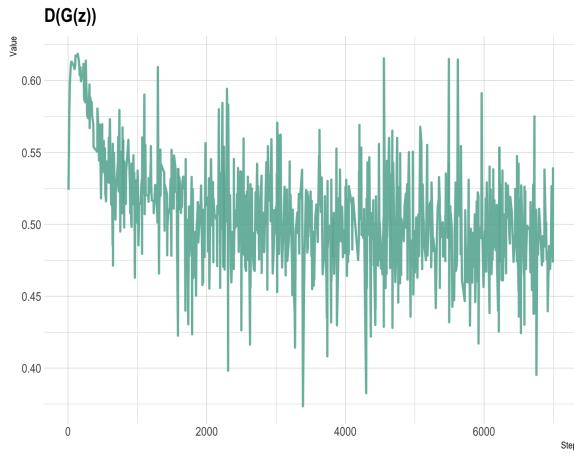
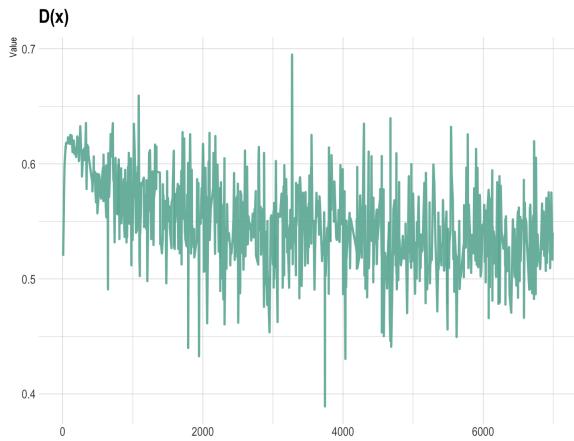
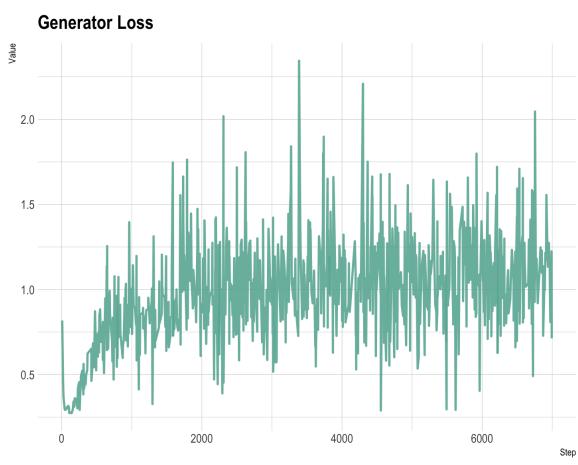
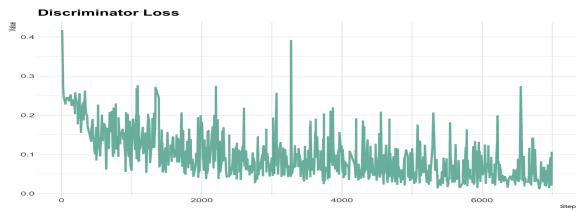


Figure 11. Generator Loss



Synthetic Pizza Dataset



For both real and synthetic pizza datasets:

We can observe that both the metrics D_x and $D(G(z))$ tend to 0.5. This means that after training for more epochs our discriminator is no longer able to discriminate between the real and fake images and hence our goal to train a generator that can fool the discriminator is realized.

Generator Loss increases as expected and the Discriminator Loss decreases as expected too.

5. Quantitative Evaluation

5.1. Introduction

The two most common GAN evaluation measures are Inception Score (IS) [5] and Fréchet Inception Distance (FID) [1]. In this section, we use FID and IS to evaluate

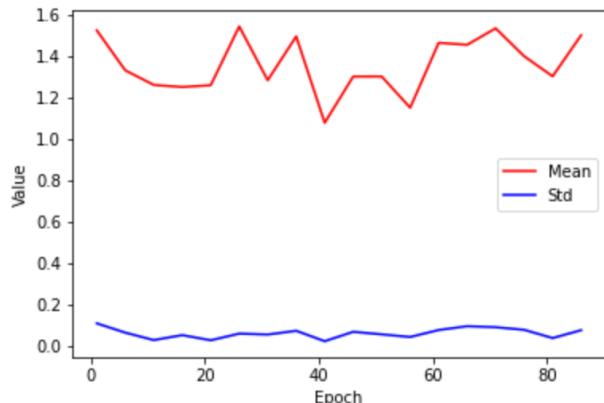
432 our GAN results quantitatively. Code Implementations see
 433 FID_IS.ipynb
 434

435 5.2. Inception Score (IS)

436 IS computes the KL divergence between the conditional
 437 class distribution and the marginal class distribution over
 438 the generated data.[\[5\]](#). IS does not capture intra-class diver-
 439 sity, is insensitive to the prior distribution over labels (hence
 440 is biased towards ImageNet dataset and Inception model4),
 441 and is very sensitive to model parameters and implemen-
 442 tations.
 443

444 5.2.1 IS for generated real pizza

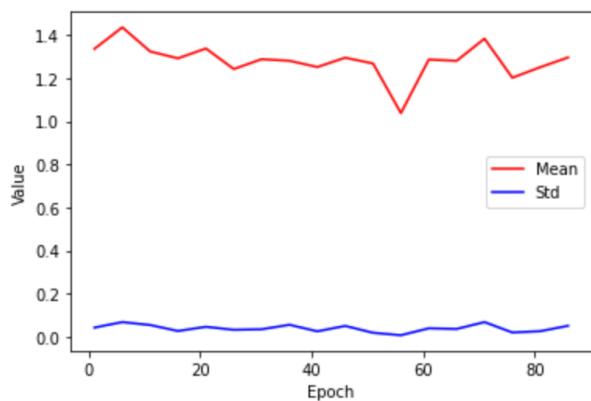
445 Figure.[18](#) shows the inception score of generated real pizza
 446 from epoch 1 to epoch 85.
 447



463 Figure 18. FID for real generated images
 464

465 5.2.2 IS for generated synthetic pizza

466 Figure.[19](#) shows the inception score of generated synthetic
 467 pizza from epoch 1 to epoch 85.
 468



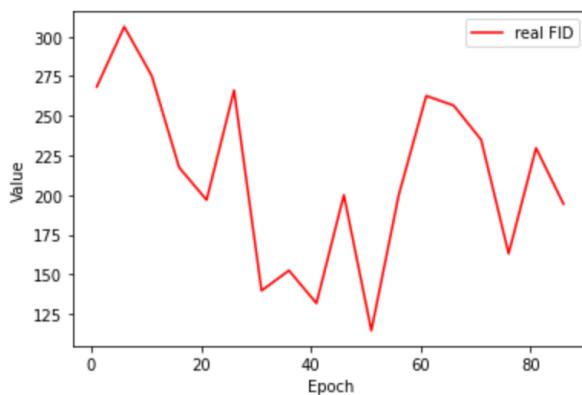
484 Figure 19. FID for synthetic generated images
 485

486 5.3. Fréchet Inception Distance (FID)

487 Unlike the earlier inception score (IS)[\[5\]](#), which evaluates
 488 only the distribution of generated images, FID [\[1\]](#) cal-
 489 culates the Wasserstein-2 (a.k.a. Fréchet) distance between
 490 multivariate Gaussians fitted to the embedding space of the
 491 Inception-v3 network of generated and real images[\[5\]](#). And it can detect intra-class mode collapse.
 492

493 5.3.1 FID for generated real pizza

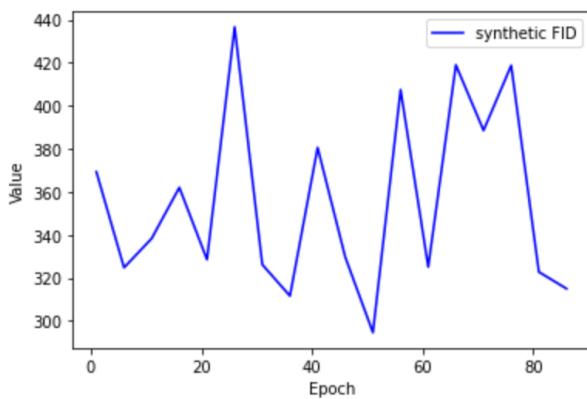
494 Figure.[20](#) shows the FID of generated real pizza from epoch
 495 1 to epoch 85.
 496



500 Figure 20. FID for real generated images
 501

502 5.3.2 FID for generated synthetic pizza

503 Figure.[21](#) shows the FID of generated synthetic pizza from
 504 epoch 1 to epoch 85.
 505



522 Figure 21. FID for synthetic generated images
 523

540

6. Analysis

541

6.1. IS for generated real pizza

543

We can observe that the inception score [5] increases slightly with increase in the number of epochs. Increase in the IS indicates better performance of the model, which is inline with our generated images.

544

545

546

547

548

549

550

551

552

553

554

555

556

557

558

559

560

561

562

563

564

565

566

567

568

569

570

571

572

573

574

575

576

577

578

579

580

581

582

583

584

585

586

587

588

589

590

591

592

593

Similar to the pizza images generated on the RealPizza Dataset, there is an increase in the IS [5] with increase in the number of epochs, showing improvement in the performance of the model. This again is consistent with our generated images.

6.3. FID for generated real pizza

For the pizza images generated on the RealPizza dataset, the FID score [1] has a value of 275 initially. With increase in the number of epochs, there is an overall decrease in the FID. This implies better performance of the model with increase in the number of epochs, which is consistent with our generated images. The FID score goes as low as 125. With increase in the number of epochs and decrease in batch size, we can further improve the performance of the model, thereby further decreasing the FID.

6.4. FID for generated synthetic pizza

For the pizza images generated on the SyntheticPizza dataset, the FID score [1] has an initial value of around 370. We can see that similar to the pizza images generated on the RealPizza dataset, there is an overall decrease in the FID score with increase in the number of epochs, indicating improvement in the performance of the model. This is inline with the quality of our generated images. The minimum value of the FID score obtained here is slightly below 300. This can be further reduced by increasing the number of epochs and decreasing the batch size.

7. Conclusion

In this project step we tackle the problem of Generating fake images using Generative Adversarial Networks. We trained two GANs, one for the real pizza images and one for the synthetic pizza images. We observed that the images generated by the real pizza GAN were very visually appealing and looked very real. We can achieve better performance by using a smaller batch size and training more epochs.

Further, we were able to generate good results with 85 epochs of training and if we were to train 100-150 epochs and augment more data either via blurring , rotating images ,adding jitters, we can achieve better results. Another future scope of work would be to fine tune the architecture to get better results.

References

- [1] M. Heusel, H. Ramsauer, T. Unterthiner, B. Nessler, and S. Hochreiter. Gans trained by a two time-scale update rule converge to a local nash equilibrium. 2017. 4, 5, 6
- [2] P. Isola, J.-Y. Zhu, T. Zhou, and A. A. Efros. “image-to-image translation with conditional adversarial networks” in *Proceedings of the IEEE conference on computer vision and pattern recognition*. pages 1125–1134, 2011. 2
- [3] J. Johnson, A. Alahi, and L. Fei-Fei. “perceptual losses for real-time style transfer and super-resolution” in *European conference on computer vision*. pages 694–711, 2016. 2
- [4] D. P. Papadopoulos, Y. Tamaazousti, F. Ofli, I. Weber, and A. Torralba. How to make a pizza: Learning a compositional layer-based gan model, Jan 1970. 1
- [5] T. Salimans, I. Goodfellow, W. Zaremba, V. Cheung, A. Radford, and X. Chen. Improved techniques for training gans, 2016. 4, 5, 6

594

595

596

597

598

599

600

601

602

603

604

605

606

607

608

609

610

611

612

613

614

615

616

617

618

619

620

621

622

623

624

625

626

627

628

629

630

631

632

633

634

635

636

637

638

639

640

641

642

643

644

645

646

647